

# Threads

POSIX ref.: <http://pubs.opengroup.org/onlinepubs/9699919799/>



SF STATE

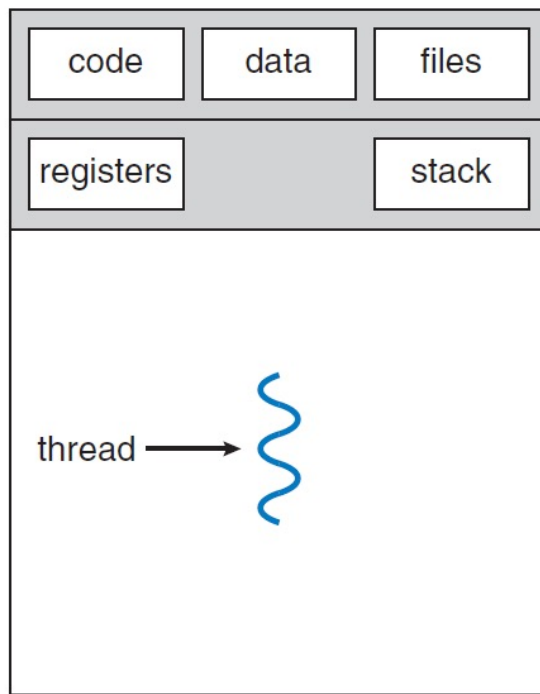
# Threads



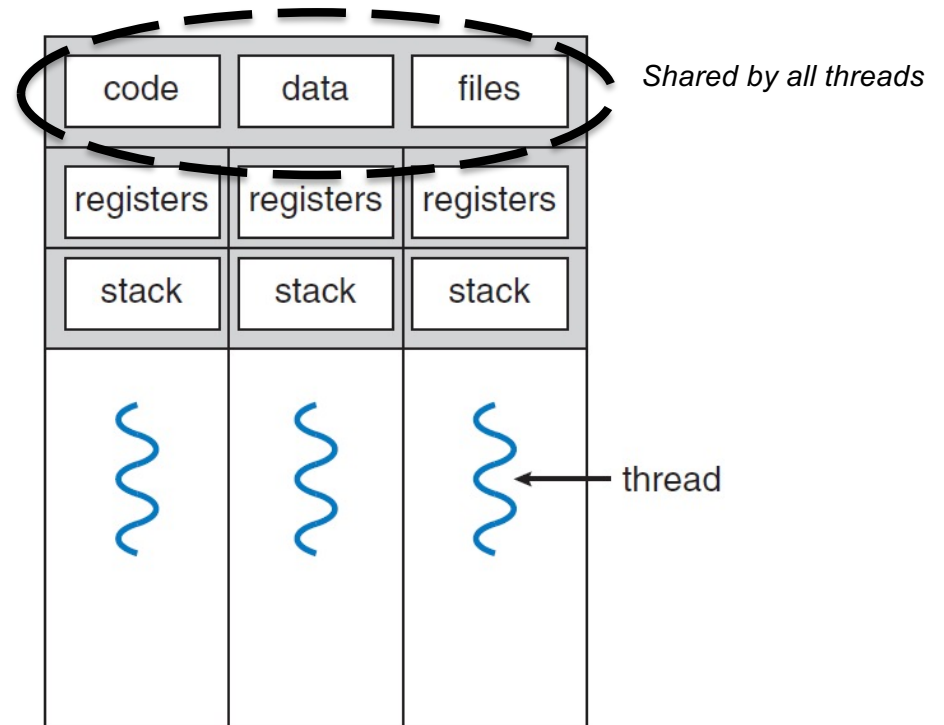
- Threads share heap, data segment, text segment and file descriptor table but have separate stacks and CPU contexts.
- Each thread have a private program counter and threads executes concurrently.
- Depending on how threads are implemented, the CPU contexts can be stored in user space or kernel space.

# Threads in processes

- A thread is a basic unit of CPU utilization



Single-threaded Process



Multithreaded Process

# Threads vs. Processes



How are threads and processes different?

...

Most modern applications are multithreaded

- E.g., word processor, Web browser, ...  
(check with `top`!)

Most OS kernels are also multithreaded

- Process creation: slow, needs more resources
- Thread creation: faster, needs fewer resources

Switching processes is also more expensive than switching threads.

# Threads/Processes speed comparison



From:

<https://computing.llnl.gov/tutorials/pthreads/#Pthread>

Time to create 50k processes/threads (in seconds):

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6GHz Xeon (16 cores/node)	8.1	0.1	2.9	0.5	0.2	0.3
Intel 2.8GHz Xeon (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
Intel 2.4GHz Xeon (2 cores/node)	54.9	1.5	20.8	1.6	0.7	0.9

# Benefits of multiple threads



Benefits of threads:


- Responsiveness
  - If current thread is waiting for IO, or busy with a slow operation
    - Can switch to other thread
- Resource sharing
  - Threads share memory and resources by default
  - Easier code than shared memory/message passing in processes
- Economy
  - Cheaper to create/switch threads vs. processes
- Scalability
  - Threads can be running in parallel in multi-processor architecture

# Multicore/Multithreaded Programming




- Concurrency:
  - Usually for single core system
  - One thread runs at a time
  - Switch threads to overlap idle time
- Parallelism:
  - For multi-core/multiprocessor system
  - Multiple threads run at a time, one on each core
  - Data Parallelism:
    - Each thread has a subset of the full data set
    - Each thread performs same operation on data subset
  - Task Parallelism:
    - Each thread performs different operation on data set

# Multicore Execution



Concurrent execution on a single-core system



Parallel execution on a multi-core system



# Non-atomic operations



- Incrementing and decrementing a variable are examples of non-atomic operations.
- The following C expression,  
`X++;`
- Is translated by the compiler to three instructions:
  - **load** X from memory into a CPU register.
  - **increment** X and save result in a CPU register.
  - **store** result back to memory.

# Race condition



- A race condition or race hazard is the behavior of system where the output is dependent on the sequence or timing of other uncontrollable events.
- It becomes a bug when events do not happen in the intended order.
- A data race occurs when two instructions from different threads access the same memory location and:
  - at least one of these accesses is a write
  - and there is no synchronization that is mandating any particular order among these accesses.

# Interleaving of executing threads



Thread A (increment)				Thread B (decrement)		
OP	Operands	\$t0	BALANCE	OP	Operands	\$t0
			0			
lw	\$t0, BALANCE	0	0			
			0	lw	\$t0, BALANCE	0
addi	\$t0, \$t0, 1	1	0			
sw	\$t0, BALANCE	1	1			
			1	addi	\$t0, \$t0, -1	-1
			-1	sw	\$t0, BALANCE	-1

# The "Too Much Milk" problem



Time	You	Your Roommate
3:00	Arrive Home	
3:05	Look in fridge, no milk	
3:10	Leave for grocery	
3:15		Arrive Home
3:20		Look in fridge, no milk
3:25		Leave for grocery
3:30	Arrive at grocery store	
3:35	Buy milk	
3:40	Arrive home, put milk in fridge	
3:45		Arrive at grocery store
3:50		Buy milk
3:55		Arrive home, put milk in fridge
4:00		Uh oh!

# Solving the Too Much Milk Problem



## Correctness properties

- Only one person buys milk
  - Safety: "nothing bad happens"
- Someone buys milk if you need to
  - Progress: "something good eventually happens"

## First: use atomic loads stores as building blocks

- "Leave a note" (lock)
- "Remove a note" (unlock)
- "Don't buy milk if there's a note" (wait)

# Too Much Milk: Attempt 1



- Thread A
  - if(no milk && no note)
    - leave note
    - buy milk
    - remove note
- Thread B
  - if(no milk && no note)
    - leave note
    - buy milk
    - remove note

# Too Much Milk: Attempt 1



- Thread A
  - if(no milk && no note)
    - leave note
    - buy milk
    - remove note
- Thread B
  - if(no milk && no note)
    - leave note
    - buy milk
    - remove note

If both threads switch after reading the if statement, then both will buy milk.

# Too Much Milk: Attempt 2



Idea: use labeled notes

- Thread A
  - leave note A
  - if (no note B)
    - if (no milk)
      - buy milk
  - remove note A
- Thread B
  - leave note B
  - if (no note A)
    - if (no milk)
      - buy milk
  - remove note B



# Too Much Milk: Attempt 2



Idea: use labeled notes

- Thread A
  - leave note A
  - if (no note B)
    - if (no milk)
      - buy milk
  - remove note A
- Thread B
  - leave note B
  - if (no note A)
    - if (no milk)
      - buy milk
  - remove note B

If both threads switch after leaving a note,  
neither will buy milk.

# Too Much Milk: Attempt 3



Idea: *wait* for the right note

- Thread A
  - leave note A
  - while (no note B)
    - do nothing
  - if (no milk)
    - buy milk
  - remove note A
- Thread B
  - leave note B
  - if (no note A)
    - if (no milk)
      - buy milk
  - remove note B

# Too Much Milk: Attempt 3



Idea: *wait* for the right note

- Thread A
  - leave note A
  - while (no note B)
    - do nothing
  - if (no milk)
    - buy milk
  - remove note A
- Thread B
  - leave note B
  - if (no note A)
    - if (no milk)
      - buy milk
  - remove note B
- Possibility 1: thread A runs first, then thread B.
  - thread A will buy milk. thread B will see there is milk, and therefore not buy any more milk.

# Too Much Milk: Attempt 3



Idea: *wait* for the right note

- Thread A
  - leave note A
  - while (no note B)
    - do nothing
  - if (no milk)
    - buy milk
  - remove note A
- Thread B
  - leave note B
  - if (no note A)
    - if (no milk)
      - buy milk
  - remove note B
- Possibility 2: thread B runs first, then thread A.
  - thread B will buy milk. thread A will see there is milk, and therefore not buy any more milk.

# Too Much Milk: Attempt 3



Idea: *wait* for the right note

- Thread A
  - leave note A
  - while (no note B)
    - do nothing
  - if (no milk)
    - buy milk
  - remove note A
- Thread B
  - leave note B
  - if (no note A)
    - if (no milk)
      - buy milk
  - remove note B
- Possibility 3: Interleaved - A waits and A buys
  - thread A and B can leave a note, but thread A will wait for thread B to finish before moving on.
  - thread B will not buy milk because note A exists, and finish the thread
  - thread A will continue and buy milk.

# Too Much Milk: Attempt 3



Idea: *wait* for the right note

- Thread A
  - leave note A
  - while (no note B)
    - do nothing
  - if (no milk)
    - buy milk
  - remove note A
- Thread B
  - leave note B
  - if (no note A)
    - if (no milk)
      - buy milk
  - remove note B
- Possibility 4: Interleaved - A waits and B buys
  - thread B will leave a note, and see there is no note A.
  - A can start as see there is a note B, and wait for B to finish.
  - B will continue to buy milk if there no milk, and then finish before thread A will continue.
  - thread A will not buy milk because thread B just bought milk.

# Critical section



- Concurrent accesses to shared resources can lead to unexpected or erroneous behavior
- Parts of the program where the shared resource is accessed are protected.
- This protected section is the **critical section** or **critical region**.
- Typically, the critical section accesses a shared resource, such as a data structure, a peripheral device, or a network connection, that would not operate correctly in the context of multiple concurrent accesses

# Mutual exclusion (mutex)



- In computer science, **mutual exclusion** is a property of **concurrency control**
- This control is instituted for the purpose of **preventing race conditions**;
  - it is the requirement that one thread of execution never enter its critical section at the same time that another concurrent thread of execution enters its own critical section.
- Often the word **mutex** is used as a short form of mutual exclusion.



# User vs Kernel Threads




- There is the concept of
  - User Threads, that is threads that are managed entirely in the user space
  - Kernel Threads, threads that are known by and managed by the Kernel
- There are advantages and disadvantages to both. But since this is a Linux Programming course, we will ignore the user threads and focus on the kernel threads.


# Threads



- All threads in a process share the Text, Data (initialized and uninitialized), heap, and library segments
- As well as
  - Open file descriptors
  - Signals
  - Current working directory
  - User and group IDs
- They do not share the Stack segment
- Nor do they share the CPU state (Program counter or Registers)
- They also have unique Thread IDs



# Linux Process/Thread Model



# Thread Example



- `pthread_t` are return structures
- `pthread_create` creates and starts a new thread

## Program Code:

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *print_message_function( void *ptr ) {
    char *message;
    message = (char *) ptr;
    printf("%s \n", message);
}

main() {
    pthread_t thread1, thread2;
    char *message1 = "Thread 1";
    char *message2 = "Thread 2";
    int iret1, iret2;

    /* Create independent threads each of which will execute function */

    iret1 = pthread_create( &thread1, NULL, print_message_function, (void*) message1);
    iret2 = pthread_create( &thread2, NULL, print_message_function, (void*) message2);
```

# Thread Example



- `pthread_join` is similar to the `wait` we used for `fork`'d processes
- But threads do NOT have a child – parent relationship any thread can wait on any other thread

## Program Code:

```
/* Wait till threads are complete before main continues. */  
/* Unless we wait we run the risk of executing an exit */  
/* which will terminate the process and all threads before */  
/* the threads have completed. */  
  
pthread_join( thread1, NULL);  
pthread_join( thread2, NULL);  
  
printf("Thread 1 returns: %d\n",iret1);  
printf("Thread 2 returns: %d\n",iret2);  
exit(0);  
}
```

## Console Output:

```
Thread 1  
Thread 2  
Thread 1 returns: 0  
Thread 2 returns: 0
```

# pthread\_create



Function call: **pthread\_create** Arguments:

```
int pthread_create(pthread_t * thread,  
                  const pthread_attr_t * attr,  
                  void * (*start_routine)(void *),  
                  void *arg);
```

- thread - returns the thread id. (unsigned long int defined in bits/pthreadtypes.h)
- attr - Set to NULL if default thread attributes are used.
- void \* (\*start\_routine) - pointer to the function to be threaded. Function has a single argument: pointer to void.
- \*arg - pointer to argument of function. To pass multiple arguments, send a pointer to a structure.

# pthread\_create attributes (1 of 2)



Attribute	Meaning	APIs: pthread_attr_[...] (3)	Values Possible	Linux Default
Detach state	Create threads as joinable or detached	pthread_attr_[get set]detachstate	PTHREAD_CREATE_JOINABLE PTHREAD_CREATE_DETACHED	PTHREAD_CREATE_JOINABLE
Scheduling/contention scope	Set of threads against which we compete for resources (CPU)	pthread_attr_[get set]scope	PTHREAD_SCOPE_SYSTEM PTHREAD_SCOPE_PROCESS	PTHREAD_SCOPE_SYSTEM
Scheduling/inheritance	Determines whether scheduling attributes are inherited implicitly from calling a thread or explicitly from the attr structure	pthread_attr_[get set]inheritsched	PTHREAD_INHERIT_SCHED PTHREAD_EXPLICIT_SCHED	PTHREAD_INHERIT_SCHED
Scheduling/policy	Determines the scheduling policy of the thread being created	pthread_attr_[get set]schedpolicy	SCHED_FIFO SCHED_RR SCHED_OTHER	SCHED_OTHER

# pthread\_create attributes (2 of 2)



Scheduling/priority	Determines the scheduling priority of the thread being created	pthread_attr_ [get set]schedparam	struct sched_param holds int sched_priority	0 (non real-time)
Stack/guard region	A guard region for the thread's stack	pthread_attr_ [get set]guardsize	Stack guard region size in bytes	1 page
Stack/location, size	Query or set the thread's stack location and size	pthread_attr_ [get set]stack pthread_attr_ [get set]stackaddr pthread_attr_ [get set]stacksize	Stack address and/or stack size, in bytes	Thread Stack Location: left to the OS Thread Stack Size: 8 MB



# Thread Example



- Here we explicitly set a thread as joinable using the attribute parameter

## Program Code:

```
main() {
    pthread_t thread1, thread2;
    pthread_attr_t attr;
    char *message1 = "Thread 1";
    char *message2 = "Thread 2";
    int iret1, iret2;

    /* Init the thread attribute structure to defaults */
    pthread_attr_init(&attr);

    /* Create all threads as joinable */
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);

    /* Create independent threads each of which will execute function */

    iret1 = pthread_create( &thread1, &attr, print_message_function, (void*) message1);
    iret2 = pthread_create( &thread2, &attr, print_message_function, (void*) message2);
}
```

# Terminating a thread



- A thread terminates when:
  - It reaches the end of its start\_routine
  - It calls pthread\_exit
  - It is canceled by another thread calling pthread\_cancel
  - The process that contains the thread terminates, for example, because of a thread calling exit, or the process receives a signal that is not handled, masked, or ignored
- Note that if a multithreaded program calls fork, only the thread that made the call will exist in the new child process. Fork does not replicate all threads.

# pthread\_exit



```
#include <pthread.h>
void pthread_exit(void *retval);
```

- The parameter specifies the exit status of the calling thread;
- This can be NULL
- A casted value
  - `pthread_exit((void *)1);`
- Or a pointer to a structure

- A join is performed when one wants to wait for a thread to finish.
- A thread calling routine may launch multiple threads then wait for them to finish to get the results.
- Unlike `fork` – any thread may wait for another thread, not just the “parent”
- `int pthread_join(pthread_t thread, void **retval);`
  - The first parameter to `pthread_join`, `thread`, is the ID of the thread to wait for.
  - The moment it terminates, the calling thread will receive, in the second parameter (a value-result style parameter), the return value from the thread that terminated—which, of course is the value passed via its `pthread_exit` call.

# Mutexes



- Mutexes are used to prevent data inconsistencies due to race conditions.
- Mutexes are used for serializing shared resources.
- Anytime a global resource is accessed by more than one thread the resource should have a Mutex associated with it.
- One can apply a mutex to protect a segment of memory ("critical region") from other threads.
- Mutexes can be applied only to threads in a single process and do not work between processes as do semaphores.

# Mutex Example



- Main Code

## Program Code:

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

void *functionC();
pthread_mutex_t mutex1 = PTHREAD_MUTEX_INITIALIZER;
int counter = 0;

main() {
    int rc1, rc2;
    pthread_t thread1, thread2;

    /* Create independent threads each of which will execute functionC */

    if ((rc1=pthread_create( &thread1, NULL, &functionC, NULL)))
    {
        printf("Thread creation failed: %d\n", rc1);
    }
```

# Mutex Example



- Main code

## Program Code:

```
if( (rc2=pthread_create( &thread2, NULL, &functionC, NULL)) )
{
    printf("Thread creation failed: %d\n", rc2);
}

/* Wait till threads are complete before main continues. Unless we */
/* wait we run the risk of executing an exit which will terminate */
/* the process and all threads before the threads have completed. */

pthread_join( thread1, NULL);
pthread_join( thread2, NULL);

exit(0);
}
```

# Mutex Example



- Thread code

## Program Code:

```
void *functionC()
{
    pthread_mutex_lock( &mutex1 );
    counter++;
    printf("Counter value: %d\n",counter);
    pthread_mutex_unlock( &mutex1 );
}
```



# Mutex Locks



- Two main functions:
  - `pthread_mutex_lock (pthread_mutex_t * mutex);`
  - `pthread_mutex_unlock (pthread_mutex_t * mutex);`
- The first start a critical section, the second ends it.

# Mutex Locks



- In order to use a mutex lock, one must first initialize it to the unlocked state; this can be done as follows:

```
if ((ret = pthread_mutex_init(&mylock, NULL)))  
    FATAL("pthread_mutex_init() failed! [%d]\n", ret);
```

- Alternatively, we could perform the initialization as a declaration, such as:

```
pthread_mutex_t mylock = PTHREAD_MUTEX_INITIALIZER;
```

- Also, once we are done, we must destroy the mutex lock(s):

```
if ((ret = pthread_mutex_destroy(&mylock)))  
    FATAL("pthread_mutex_destroy() failed! [%d]\n", ret);
```

# Wait on first to finish



- How to wait on finish order, rather than creation order

## Program Code:

```
void *routine(void *arg)
{
    int *fds = (int *)arg;

    pthread_t t = pthread_self(); /* Get Thread ID */

    usleep((rand()/(1.0 + RAND_MAX)) * 1000000);

    write(fds[1], &t, sizeof(t));
}
```

# Wait on first to finish



- How to wait on finish order, rather than creation order

## Program Code:

```
int main() {
    int fds[2];
    srand(time(0));
    pipe(fds);
    for (int i = 0; i < 2; i++) {
        pthread_t tid;
        pthread_create(&tid, NULL, routine, fds);
        printf("created: %llu\n", (unsigned long long)tid);
    }
    for (int i = 0; i < 2; i++) {
        pthread_t tid;
        read(fds[0], &tid, sizeof(tid)); /*Block on Pipe, 1st to write, 1st done*/
        printf("joining: %llu\n", (unsigned long long)tid);
        pthread_join(tid, 0);
    }
}
```

# Daemons



- A daemon is a process that runs in the background, owned by the init process and not connected to a controlling Terminal.
- The steps to create a daemon are as follows:
  1. Call `fork` to create a new process, after which the parent should exit, thus creating an orphan which will be re-parented to `init`.
  2. The child process calls `setsid`, creating a new session and process group of which it is the sole member.
    - Consider this a way of isolating the process from any controlling terminal.
  3. Change the working directory to the root directory.
  4. Close all file descriptors and redirect `stdin`, `stdout`, and `stderr` (descriptors 0, 1, and 2) to `/dev/null` so that there is no input and all output is hidden.
- Thankfully, all of the steps can be achieved with a single function call, `daemon`